

The main function contains calls to `exit()` (line 66) and `pthread_exit()` (line 80). How will the effect of these two calls differ when they are executed?

The regular `exit` will remove all the resources that were created by `main`. This includes any shared memory between processes, and is generally not advisable when you're working with threads. `Pthread_exit()` leaves process-shared resources (eg mutexes), and after the last thread in a process terminates, it will call regular `exit` with a 0 status, so that the regular `exit` routines can be processed.

The main function calls `pthread_join()` (line 77) with the parameter `thread_return`. Where does the value stored in `thread_return` come from when the `consumer_thread` is joined?

The value stored in `thread_return` is assigned by the `exit` function to a field within the `p_thread` type structure for the thread we're exiting from. That structure is keyed to a PID, so we will be able to find it that way (or if the process asking for the join is the parent by using the variable), and pull out the return value.

Where does the value stored in `thread_return` come from if the joined thread terminated by calling `pthread_exit` instead of finishing normally? On the same call to `pthread_join()` (line 77), what will it do if the thread being joined (`consumer_thread`, in this case) finishes before the main thread reaches the that line of code (line 77)?

Because the value is placed in a thread structure created elsewhere, that structure will still contain the return value of that thread. Therefore, when the code reaches the end, if `consumer` has already finished, that value will be sitting within the structure waiting.

In this program, the main thread calls `pthread_join()` on the threads it created. Could a different thread call `pthread_join()` on those threads instead? Could a thread call `pthread_join()` on the main thread (assuming it knew the main thread's thread ID - i.e. `pthread_t`)?

Yes, it could. The function waits on the finishing of a thread, which can be identified by a thread created in the context of that function, or a pointer to a thread from another process. Either way, when that thread is determined to be finished, the `pthread_exit` routine will place a return value in the `pthread_join` call. This is possible with any thread, including `main`.

The `consumer_routine` function calls `sched_yield()` (line 180) when there are no items in the queue. Why does it call `sched_yield()` instead of just continuing to check the queue for an item until one arrives?

`Sched_yield()` allows other threads to take CPU time. If we just kept attempting to grab the lock, then one thread could feasibly retain all the CPU uptime and the other threads would never process.